

The results are in! See what nearly 90,000 developers picked as their most loved, dreaded, and desired coding languages and more in the 2019 Developer Survey.

\d is less efficient than [0-9]

[Ask Question](#)

1179



232

I made a comment yesterday on an answer where someone had used `[0123456789]` in a [regular expression](#) rather than `[0-9]` or `\d`. I said it was probably more efficient to use a range or digit specifier than a character set.

I decided to test that out today and found out to my surprise that (in the C# regex engine at least) `\d` appears to be less efficient than either of the other two which don't seem to differ much. Here is my test output over 10000 random strings of 1000 random characters with 5077 actually containing a digit:

```
Regular expression \d          took 00:00:00.2141226 result: 5077/10000
Regular expression [0-9]      took 00:00:00.1357972 result: 5077/10000 63.42 % of
first
Regular expression [0123456789] took 00:00:00.1388997 result: 5077/10000 64.87 % of
first
```

It's a surprise to me for two reasons:

1. I would have thought the range would be implemented much more efficiently than the set.
2. I can't understand why `\d` is worse than `[0-9]`. Is there more to `\d` than simply shorthand for `[0-9]`?

Here is the test code:

```

using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;
using System.Diagnostics;
using System.Text.RegularExpressions;

namespace SO_RegexPerformance
{
    class Program
    {
        static void Main(string[] args)
        {
            var rand = new Random(1234);
            var strings = new List<string>();
            //10K random strings
            for (var i = 0; i < 10000; i++)
            {
                //Generate random string
                var sb = new StringBuilder();
                for (var c = 0; c < 1000; c++)
                {
                    //Add a-z randomly
                    sb.Append((char)('a' + rand.Next(26)));
                }
                //In roughly 50% of them, put a digit
                if (rand.Next(2) == 0)
                {
                    //Replace one character with a digit, 0-9
                    sb[rand.Next(sb.Length)] = (char)('0' + rand.Next(10));
                }
                strings.Add(sb.ToString());
            }

            var baseTime = testPerformance(strings, @"\d");
            Console.WriteLine();
            var testTime = testPerformance(strings, "[0-9]");
            Console.WriteLine(" {0:P2} of first", testTime.TotalMilliseconds /
baseTime.TotalMilliseconds);
            testTime = testPerformance(strings, "[0123456789]");
            Console.WriteLine(" {0:P2} of first", testTime.TotalMilliseconds /
baseTime.TotalMilliseconds);
        }

        private static TimeSpan testPerformance(List<string> strings, string regex)
        {

```

```

var sw = new Stopwatch();

int successes = 0;

var rex = new Regex(regex);

sw.Start();
foreach (var str in strings)
{
    if (rex.Match(str).Success)
    {
        successes++;
    }
}
sw.Stop();

Console.WriteLine("Regex {0,-12} took {1} result: {2}/{3}", regex, sw.Elapsed,
successes, strings.Count);

return sw.Elapsed;
}
}
}

```

c#

regex

performance

edited May 22 '13 at 21:06



Peter Mortensen

13.9k 19 87 113

asked May 18 '13 at 7:18



weston

39.7k 16 96 171

167 Maybe \d deals with locales. E.g. Hebrew uses letters for digits. – Barmar May 18 '13 at 7:20

6 related: stackoverflow.com/a/6479605/674039 – wim May 18 '13 at 15:04

35 This is an interesting question precisely because \d does not mean the same thing in different languages. In Java, for example [\d does](#)

Home

PUBLIC


Stack Overflow

Tags

Users

Q&A for work

[Learn More](#)

7 Porting weston's code to Java yields: -- Regex \d took
00:00:00.043922 result: 4912/10000 -- Regex [0-9] took
00:00:00.073658 result: 4912/10000 167% of first -- Regex
[0123456789] took 00:00:00.085799 result: 4912/10000 195% of first --
[Lunchbox](#) May 22 '13 at 16:35 

6 Answers



`\d` checks all Unicode digits, while `[0-9]` is limited to these 10 characters. For example, [Persian](#) digits, `۱۲۳۴۵۶۷۸۹`, are an example of Unicode digits which are matched with `\d`, but not `[0-9]`.

You can generate a list of all such characters using the following code:

```
var sb = new StringBuilder();
for (UInt16 i = 0; i < UInt16.MaxValue; i++)
{
    string str = Convert.ToChar(i).ToString();
    if (Regex.IsMatch(str, @"\d"))
        sb.Append(str);
}
Console.WriteLine(sb.ToString());
```

Which generates:

0123456789၁၂၃၄၅၆၇၈၉၀.၁၂၃၄၅၆၇၈၉.၁၂၃၄၅၆၇၈၉၀၁၂၃၄၅၆၇၈၉၀၁
၁၂၃၄၅၆၇၈၉၀၁၂၃၄၅၆၇၈၉၀၁၂၃၄၅၆၇၈၉၀၁၂၃၄၅၆၇၈၉၀၁၂၃၄၅၆၇၈၉၀၁၂

edited Oct 13 '16 at 17:35



3,419 7 28 48



13.1k 3 22 42

- <https://stackoverflow.com/questions/16621738/d-is-less-efficient-than-0-9>

17

An addition to [top answer](#) from [Sina Iravanian](#), here is a .NET 4.5 version (since only that version supports UTF16 output, c.f. the first three lines) of his code, using the full range of Unicode code points. Due to the lack of proper support for higher Unicode planes, many people are not aware of always checking for and including the upper Unicode planes. Nevertheless they sometimes do contain some important characters.

Update

Since `\d` does not support non-BMP characters in regex (thanks [xanatos](#)), here a version that uses the Unicode character database

```
public static void Main()
{
    var unicodeEncoding = new UnicodeEncoding(!BitConverter.IsLittle);
    Console.InputEncoding = unicodeEncoding;
    Console.OutputEncoding = unicodeEncoding;

    var sb = new StringBuilder();
    for (var codePoint = 0; codePoint <= 0x10ffff; codePoint++)
    {
        var isSurrogateCodePoint = codePoint <= UInt16.MaxValue
            && ( char.IsLowSurrogate((char) codePoint)
                || char.IsHighSurrogate((char) codePoint)
            );

        if (isSurrogateCodePoint)
            continue;

        var codePointString = char.ConvertFromUtf32(codePoint);

        foreach (var category in new []{
            UnicodeCategory.DecimalDigitNumber,
            UnicodeCategory.LetterNumber,
            UnicodeCategory.OtherNumber})
        {
            sb.AppendLine($"{category}");
            foreach (var ch in charInfo[category])
            {
                sb.Append(ch);
            }
            sb.AppendLine();
        }
    }
}
```

```
    }  
    Console.WriteLine(sb.ToString());  
  
    Console.ReadKey();  
}
```

Yielding the following output:

DecimalDigitNumber

[illegible]

LetterNumber

[illegible]

OtherNumber

[illegible]

edited Apr 12 '18 at 8:33



2,630 4 27 45

4 If I remember correctly, sadly in .NET Regex doesn't support non-BMP characters. So in the end checking for characters > 0xffff with a regex is useless. – [xanatos](#) Apr 12 '17 at 10:19

8/11

comparison.

-16

For example, if I wanted Regex to find IP addresses, I would rather use `\d` than `[0123456789]` or even `[0-9]` to represent any digit.

Generally speaking in my Regex use, function is more important than speed.

edited Jul 22 '17 at 20:35



[bobble bubble](#)

6,446 1 14 29

answered Feb 8 '15 at 20:56



[Faizal](#)

5

- 1 While this may incur a small penalty, it can be done once before comparing the pattern to any input string. So, the time-complexity is $O(1)$, not $O(n)$, (where n is the length of the input string.) In short, the impact is minimal at best. – [jpaugh](#) Mar 6 '18 at 18:28

`\d` checks all Unicode, while `[0-9]` is limited to these 10 characters. If just 10 digits, you should use. Others I recommend using `\d`, Because writing less.

0

answered Mar 11 '16 at 10:27



[dengkai](#)

1 2

From [Does “\d” in regex mean a digit?](#):

110

`[0-9]` isn't equivalent to `\d`. `[0-9]` matches only `0123456789` characters, while `\d` matches `[0-9]` and other digit characters, for example Eastern Arabic numerals `٠١٢٣٤٥٦٧٨٩`

edited May 23 '17 at 11:47



Community ♦

1 1

answered May 18 '13 at 7:27



Ismet Alkan

4,355 3 31 63

48 According to: msdn.microsoft.com/en-us/library/20bw873z.aspx If ECMAScript-compliant behavior is specified, \d is equivalent to [0-9]. – User 12345678 May 18 '13 at 7:30 ✎

2 huh, am i wrong or this sentence from the link is telling the opposite. "\d matches any decimal digit. It is equivalent to the \p{Nd} regular expression pattern, which includes the standard decimal digits 0-9 as well as the decimal digits of a number of other character sets." – Ismet Alkan May 18 '13 at 7:51

3 @ByteBlast thanks, using the constructor: `var rex = new Regex(regex, RegexOptions.ECMAScript);` makes them all pretty much indistinguishable in performance terms. – weston May 18 '13 at 7:53

2 oh anyway, thanks everyone. this question turned out to be a great learning for me. – Ismet Alkan May 18 '13 at 7:54

3 Please don't "just copy" answers from other questions. If the question is a duplicate, flag it as such. – BoltClock ♦ May 18 '13 at 12:00

▲ Credit to ByteBlast for noticing this in the docs. Just changing the regex constructor:

260

▼ `var rex = new Regex(regex, RegexOptions.ECMAScript);`

Gives new timings:

```
Regex \d          took 00:00:00.1355787 result: 5077/10000
Regex [0-9]       took 00:00:00.1360403 result: 5077/10000 100.34 %
```

Regex [0123456789] took 00:00:00.1362112 result: 5077/10000 100.47 %

answered May 18 '13 at 9:37



weston

39.7k 16 96 171

-
- 10 What does the `RegexOptions.ECMAScript` do? – laurent May 20 '13 at 1:36
-
- 5 From [Regular Expression Options](#): "Enable ECMAScript-compliant behavior for the expression." – chrisaycock May 20 '13 at 1:58
-
- 83 Effectively, I think it removes support for Unicode. – 0xFE May 20 '13 at 3:33
-
- 26 @0xFE: Not quite. Unicode escapes are still valid in ECMAScript (`\u1234`). It's "just" the shorthand character classes that change meaning (like `\d`) and the Unicode property/script shorthands that go away (like `\p{N}`). – Tim Pietzcker May 20 '13 at 9:51
-
- 9 This is not an answer to the "why" part. It is a "fix the symptoms" answer. Still valuable information. – usr May 29 '13 at 16:52 ✎
-